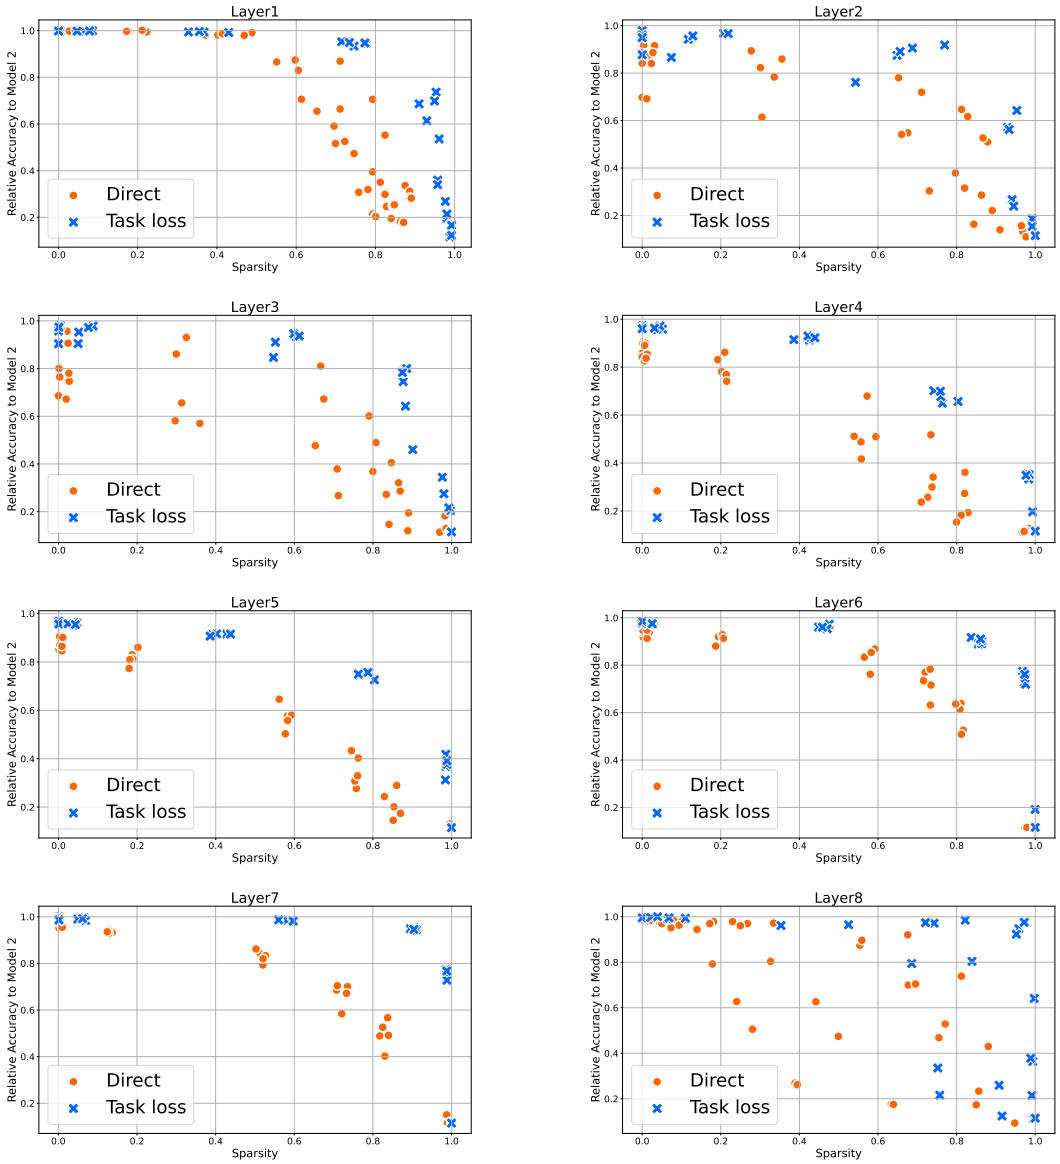

Figure 13: Relative accuracy to Model 2 with respect to sparsity. Different plots correspond to matchings on different layers.

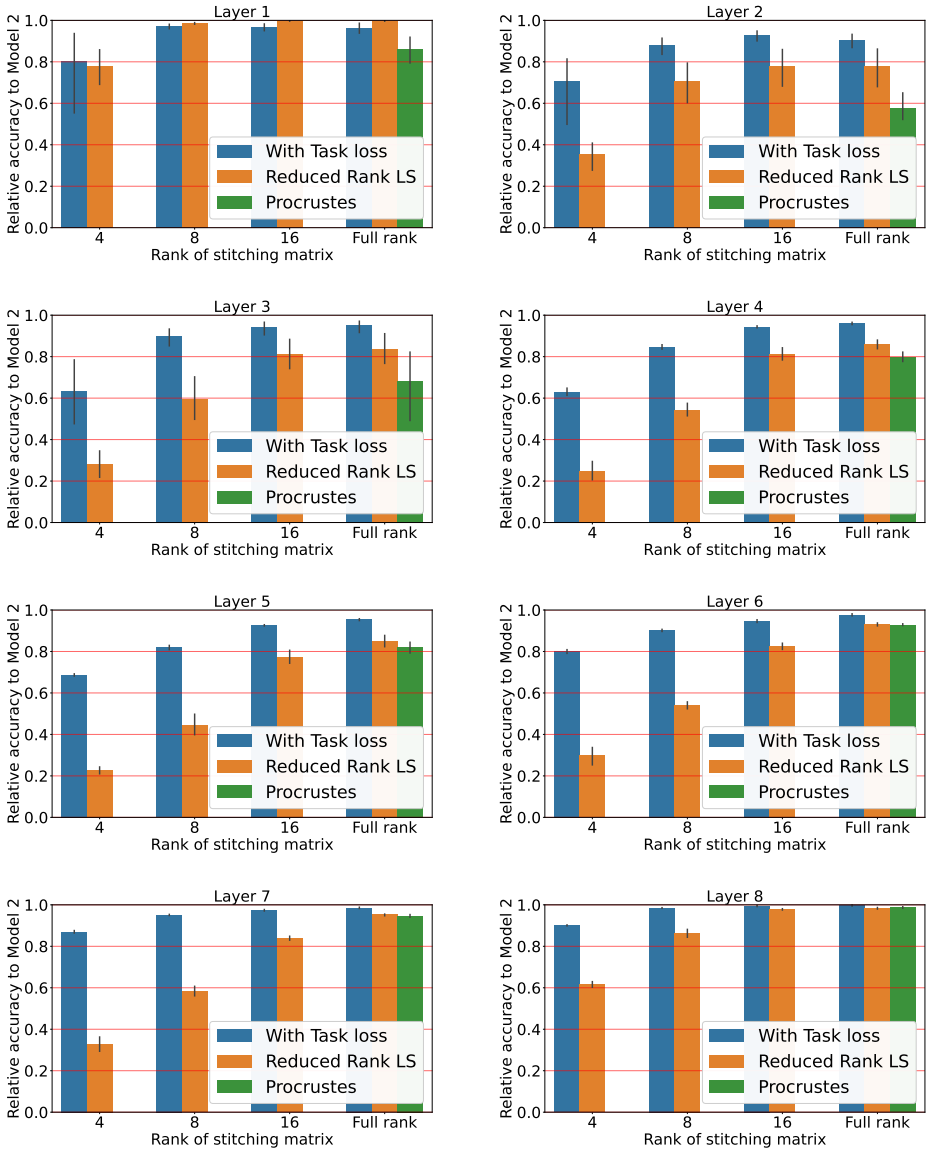

Figure 14: Performance of the low-rank analogues of least squares and task loss matchings in terms of relative accuracy. Averages of 5 runs, error bars denote standard deviations.

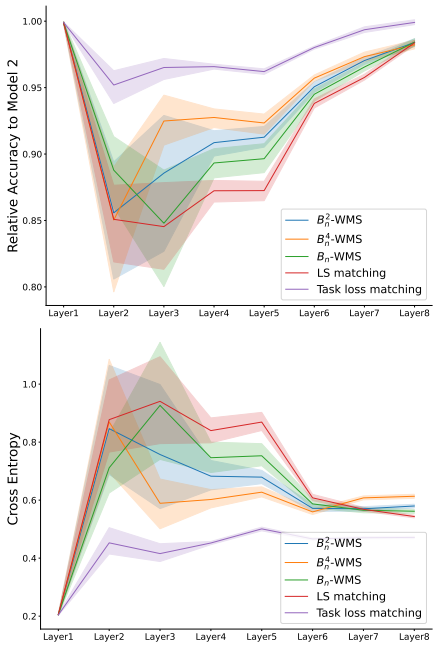

(a) Activation weighted mean squared direct matching (6), activations are raised on different exponents after normalization.

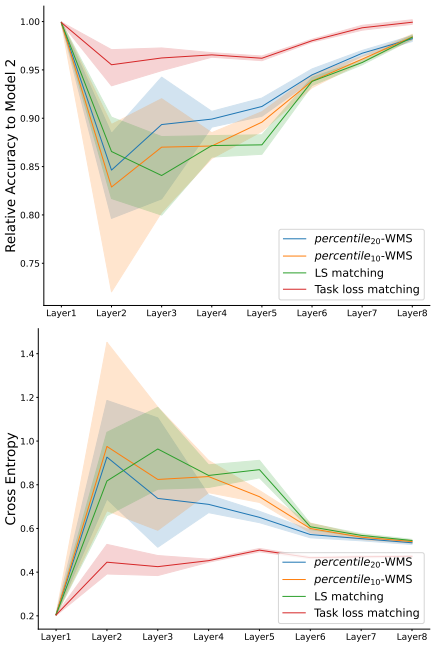

(b) 0-1 weighted mean squared direct matching (7), with the threshold T_i set according to a percentile of the activations b_{ij}, and activations below the threshold are not matched.

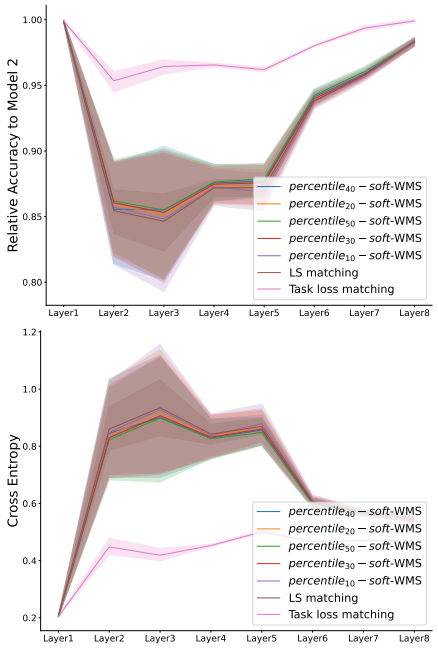

(c) Activation weighted mean squared direct matching (8), with the threshold T_i set according to a percentile of the activations b_{ij}, and activations below threshold matched to stay below threshold.

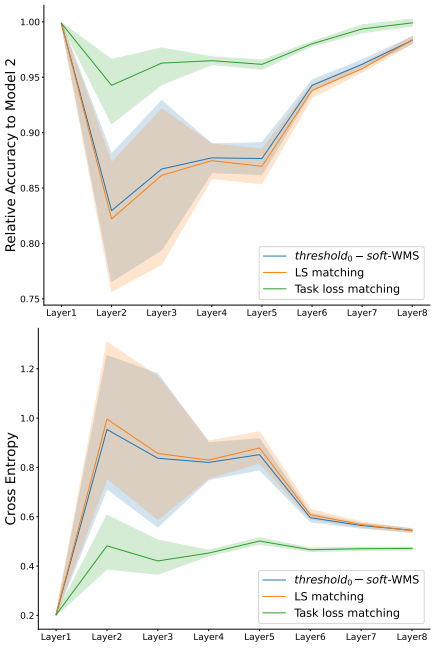

(d) 0-1 weighted mean squared direct matching (8), with the threshold T set to 0, and activations below threshold matched to stay below threshold.

Figure 15: Relative accuracy to Model 2 and cross-entropy of Tiny10's stitched layers with different weighted mean squared direct matching methods.